Independently Verifiable Stablecoin Liquidity Investigation Agent

Evidence-grounded investigation and independent replay of Ethereum Curve 3pool

WANG Hanyang HE Yuequn WU Yuxuan

5 November 2026

Abstract

This study tests whether evidence guidance and independent recomputation improve stablecoin investigation reports. We froze 51 block-pinned snapshots covering two Ethereum Curve 3pool windows and defined six tasks with 18 required components. Sixteen were answerable. We compared a deterministic template with ordinary and evidence-guided agents using the same DeepSeek Flash model, task data, tool access and budget limits.

Across 18 scheduled runs per model condition, the ordinary agent achieved reproducible correct coverage of 34/48 (70.8%) and the evidence-guided agent achieved 43/48 (89.6%). Verification withheld six nonconforming structured numeric submissions, one with an incorrect value and five with incorrect query bounds or fields. Published correct coverage was 31/48 and 42/48, respectively. On a separate 32-case test, the second-source verifier rejected 10 of 12 injected faults and returned UNKNOWN on the other two.

Research questions

Generation. Does evidence guidance change raw correctness, reproducibility or model usage under matched settings?

Verification. Which claims survive a fixed verification and withdrawal rule, and how much correct coverage is lost?

Source retrieval. What does a second archive interface add over independent arithmetic using the original snapshot?

Investigation workflow

01 / Acquire	02 / Generate	03 / Recompute	04 / Publish
Resolve blocks at UTC targets and store raw responses.	Emit required answer slots, citations and a tool trace.	Check original-source and second-source evidence separately.	Retain PASS claims and record each withdrawal and reason.

The evidence package also contains a separate local-chain digest registration demonstration. [1, 2, 3]

Data and frozen benchmark

The Ethereum Curve 3pool holds DAI, USDC and USDT. Nominal balances use the tokens' ERC-20 balanceOf(pool) values. Each target maps to a block whose timestamp is at or before the target and whose next block is later. The logs preserve target and observed times, block identities, request parameters and raw responses. The primary archive used Alchemy and the comparison archive used dRPC. The source check establishes agreement across those interfaces for the captured observations. Their upstream infrastructure is unverified. [2]

Window and UTC sampling	Snapshots	Tasks
USDC depeg context, 10 Mar 2023 00:00 to 11 Mar 00:00 inclusive, plus 11 Mar 01:00 for the dependent quote	26	R01-R03
FTX context, 9 Nov 2022 00:00 to 10 Nov 00:00 inclusive. The dependent quote falls inside this grid.	25	R04-R06

Deterministic metrics

Normalize each integer balance by dividing by 10 raised to its decimals. A token's share is its fraction of the three normalized balances. The share change is the final percentage minus the initial percentage. The quote rate is normalized output divided by input. Size sensitivity is the signed difference between the 1,000-input rate and the 100,000-input rate, divided by the 1,000-input rate and multiplied by 10,000. Imbalance is half the sum of the absolute differences between each share fraction and one third. [2]

Task family	Required components	Boundary
R01 / R04	DAI share, 1,000 USDC-to-USDT quote rate, and relative size sensitivity	Fixed hour, direction and units
R02 / R05	USDC share change, sampled maximum imbalance, and quote one hour after the returned maximum	Maximum over allowed grid points
R03 / R06	First-hour USDC share, largest allowed one-hour increase, and whether continuous peak is identifiable	Increases checked from 18:00 to 00:00 UTC, skipping the excluded 22:00 point

Controls and denominator

T is a single deterministic report per task. G and E each ran six tasks three times in a recorded randomized order. E adds requirements for precise citations, checks on units and times, and limits on interpretation. The comparison evaluates this prompt as a whole. The 16 answerable slots contribute 48 opportunities to each model arm, including failed runs. Each complete task set also has two non-answerable slots. [2, 3]

The ten exploratory rounds found a contrast in citation replayability and no clear difference in unsupported statements on the local pilot model. Those observations informed the choice of reproducible correct coverage before the formal G/E freeze. The metric rewards correct answers with complete attribution, which E explicitly requests. We also report raw correctness and the descriptive prose audit. [3, 8, 10]

The submitted workbook was attested as a consolidated human review. Its 18 components matched the frozen candidate reference and its 51 raw snapshots matched the archive. Evidence of separate A/B submissions was unavailable. [2, 4]

02 / FORMAL EXPERIMENT

Correct coverage and cost

Correct coverage requires the correct value, unit, time and query for the frozen slot. Reproducible correct coverage also requires a V1 PASS, including the exact evidence ID set, formula version and dataset digest. Publication uses the original model draws and the same denominator. [3, 6]

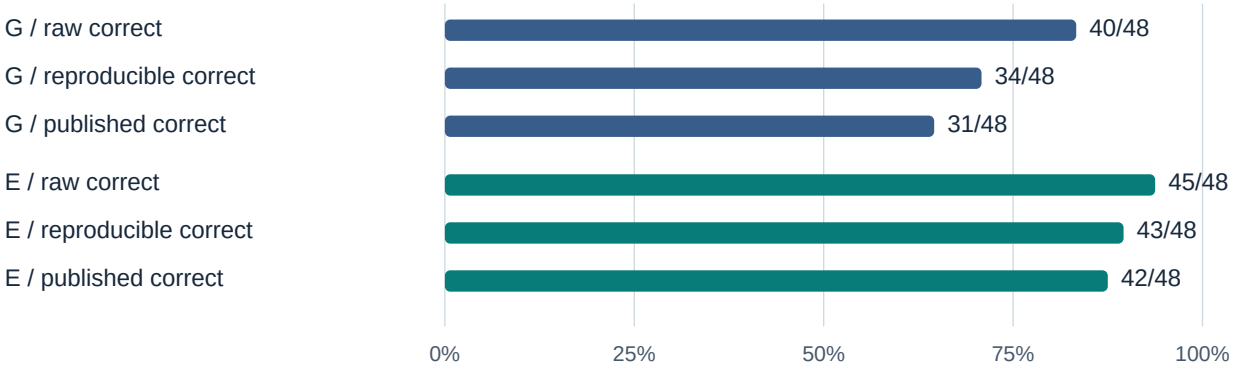

Measure	T template	G ordinary	E evidence
Scheduled runs / reports submitted	6 / 6	18 / 16	18 / 18
Correct coverage, raw	16/16 (100.0%)	40/48 (83.3%)	45/48 (93.8%)
Reproducible correct coverage, raw	16/16 (100.0%)	34/48 (70.8%)	43/48 (89.6%)
Correct coverage after publication	N/A	31/48 (64.6%)	42/48 (87.5%)
Withdrawn numeric items: value / query errors	0 / 0	1 / 2	0 / 3
Correct raw answers lost on publication	N/A	9/40 (22.5%)	3/45 (6.7%)
Over-refusals on answerable slots	0/16	0/48	0/48
Refusals passing the exact-reason check	2/2	0/6	0/6
Model calls / logged tokens	0 / 0	63 / 104,166	70 / 129,007

G's two stops omitted five answerable slots and one non-answerable slot. Excluding the two pairs with a G budget stop gives raw correctness of 40/43 in both arms and reproducible correctness of 34/43 for G and 39/43 for E. E scored four reproducible answers in the excluded pairs. The observed gap is 18.75 percentage points overall and 11.63 in this descriptive subset. [3]

The frozen refusal scorer required the literal reason code. The prompt asked for a specific reason in words. By meaning, G gave five valid refusals across six scheduled non-answerable slots, with one omitted at a budget stop. E gave six valid refusals across six slots. [3, 8]

Matched limits and actual use

Client admission used the same ceilings of 48,000 input, 12,000 output and 60,000 total tokens, with 10 calls and 2,048 output tokens per call. All 133 provider receipts were within the reserved bounds. Archived generation wall times totaled 191.602 s for G and 234.587 s for E. [3, 11]

After the study, five offline batches on Mac arm64 with Python 3.9.6 checked the archived verdicts for 102 claims across 36 preloaded records. Separate median wall times were 0.355 s for V1 and 0.435 s for V2, with zero network calls. Timings exclude generation, acquisition and publication. Fees and end-to-end costs were unmeasured. [11]

Where the checks succeed and fail

V1 recomputes claims using the frozen primary archive. V2 tests the same claim against a second source, checking the allowed observations and block identity before independently recomputing the value. A shared verifier path performs the two checks. Its arithmetic does not use the agent's intermediate results. Each version checks units, scope, evidence and numerical tolerance before publication. [5, 6]

Case set	Cases	V1 PASS / FAIL / UNKNOWN	V2 PASS / FAIL / UNKNOWN
Injected errors	12	5 / 7 / 0	0 / 10 / 2
Clean twins	12	12 / 0 / 0	12 / 0 / 0
Legitimate refusal scenarios	4	0 / 0 / 4	0 / 0 / 4
Missing secondary source	2	2 / 0 / 0	0 / 0 / 2
Unresolved source conflict	2	2 / 0 / 0	0 / 0 / 2

V2 rejected 10/12 injected faults. The three additional detections involved block identity. The other two injected faults returned UNKNOWN because source evidence was missing. Both layers accepted all 12 clean twins. The design schedules 136 case or record checks across V1 and V2, including the two empty model records in both layers. This count is separate from claim verdicts and RPC traffic. [5]

Verification and withdrawal

V1 and V2 gave identical verdicts on all 102 emitted structured claims. The verdict totals were 34 PASS, 3 FAIL and 11 UNKNOWN for G and 43 PASS, 2 FAIL and 9 UNKNOWN for E. The frozen replay withheld FAIL and UNKNOWN items and withheld PASS items when a declared dependency had been withdrawn. The publisher discarded the entire original free-text answer and rendered a new answer from surviving numeric fields. The experiment measures structured withdrawals. It provides no estimate of selective removal of unsupported prose. [3, 5, 6]

Correcting the slot-to-claim mapping restored three G and four E items. The source-conflict correction changed two scripted verdicts to UNKNOWN and preserved the formal outcomes. All frozen reports with the dependent slot declared its required parent. The current public JSON includes only verified claim fields. Formal replay compares those fields against the archive, which retains the full original reports. [3, 5, 6]

Evidence integrity on a local chain

The local Ganache demonstration registers the evidence manifest digest and reads it back. It rejects duplicate registration and produces a different, unregistered digest for a modified manifest. The saved receipt has chain ID 1337 and block number 2. Checking covered artifacts requires comparing their bytes with the manifest hashes. The arithmetic is checked by the separate verifier. [7]

What the evidence supports

Prose evidence audit

Of the 42 scheduled records, 34 contained model prose. Six deterministic template stubs and two G budget stops are outside the prose denominators. The audit recomputes numerical bundles from the frozen data and separately checks coverage by the original structured-slot evidence IDs. It also records judgments on additional material assertions and scope-limited refusals. [8]

Prose evidence measure	G ordinary	E evidence
Numerical bundles matching frozen data	42/43	48/48
Numerical bundles with incomplete original citations	6/43	2/48
Supported scope-limited refusals observed	5/5	6/6
Submitted reports with any unsupported assertion	11/16	6/18
Additional assertions, supported / unsupported / ambiguous	0 / 15 / 8	4 / 7 / 4

The counts come from one documented evidence review. The unsupported-assertion comparison is descriptive and has no independent human agreement or adjudication record. Structured coverage on page 3 uses the required answer slots, including those lost to budget stops. [3, 8]

A worked example

For R02, the allowed endpoint snapshots put USDC at 29.679648% and 48.100138% of the pool's nominal balances. The difference is +18.420491 percentage points. The sampled maximum imbalance is 0.277696115 at 11 March 2023 00:00 UTC. The 100,000 USDC-to-USDT quote rate at the required following hour is 0.9897830738 USDT/USDC. The frozen packet identifies endpoint blocks 16794061 and 16801143. [2, 8]

P41 reports a decrease of about 18.42 percentage points in both the structured item and the prose. Both verifiers reject the sign error, and the item is withheld. The endpoint balances show a share increase. Attributing that change to capital flows, solvency or a specific mechanism requires additional evidence. [2, 3, 8]

Scope of the findings

The study covers one pool, two event windows and six tasks with repeated runs. Scheduled slots, repeated agent runs and verifier calls are different units of analysis. The tested second-source benefit applies to the archived interfaces and scripted faults. Live provider availability and broader pool coverage remain open questions. [2, 3, 5]

05 / REPRODUCIBILITY

Replay and supporting records

Replay from the submission root

Action	Command
Replay deterministic data and verification	`python3 -B docs/reproduce_current.py`
Check saved formal V2 and publication outcomes	`python3 -B docs/verify_formal_replay.py`
Check packaged file hashes	`python3 -B docs/verify_delivery.py .`
Run the Python test suite	`python3 -B -m unittest discover -s tests -v`

Python 3.10 or later is recommended. The saved demonstration reproduced the two-source snapshot data and all 16 answerable components without network access or model calls. The package includes the frozen tasks, primary and comparison data, references, raw model outputs, V1/V2 decisions and corrected publication replay. See the root README for the optional local-chain replay and package-specific instructions. [2, 3, 5, 7]

Development optimization after the formal run

The compact read-window response preserves raw balances, quotes and block identifiers. Three recorded budget stops passed a counterfactual pre-dispatch admission check with the smaller payload. In two nine-run development samples using the same tasks and budget, completed submissions rose from 8/9 to 9/9 and measured token use fell from 66,475 to 62,078. The samples used different model draws and the improvement is reported separately from the formal comparison. [9]

Evidence index

[1] Course Trade 2 scope and instructor proposal feedback, 14 October 2026.

[2] results/current_preparation/{binding, dataset, secondary_dataset, tasks}.json and data/current_study/ source records.

[3] experiments/formal_v1/{freeze, result_summary, publication_replay}.json and saved raw report traces.

[4] results/current_preparation/reference_review/review_attestation.json and consolidated workbook.

[5] experiments/formal_v1/fault_batch/summary.json and experiments/formal_v1/verifier_replay_20261105.json.

[6] curve_audit/current_verify.py and curve_audit/verification.py.

[7] contracts/anchor.mjs and results/current_chain/anchor.json.

[8] experiments/formal_v1/prose_review_input.json and audit/{prose_audit_summary, prose_audit_rows}.json.

[9] results/optimization_20261105/{summary, budget_replay}.json.

[10] references/trial-report-20261102.md, exploratory trial and metric selection rationale.

[11] results/submission_validation/verifier_timing.json, generation totals and offline timings.

Freeze SHA-256 `ec66a6aaf06a0a37eeff5e74bbda6e68a866be00617a00a9413087679e7cd220`